CS 5/7320 Artificial Intelligence

More Important AI Topics

Slides by Michael Hahsler

This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International License.

Reinforcement Learning

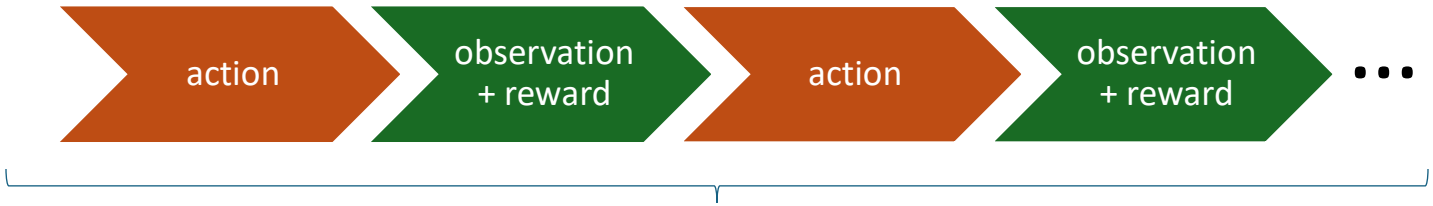

Sequential decision making: Find a policy π that maximizes the expected discounted sum of rewards over time.

$$U = \mathbb{E}\left[\sum_{t=1}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1})\right]$$

Models for the environment and the reward are known (and states evolve Markovian)
- Markov Decision Model (MDP)
- Partially Observable Markov Decision Model (POMDP)

Dynamic Programming
- Value iteration $V(s)$
- Policy iteration $\pi(s)$

Model-free approaches
- Q-Learning (learns the value of actions in states $Q(s, a)$)
- Time differencing (TD learning)

Learn iteratively

Natural Language Processing

- Speech recognition
- Informa
- Langua

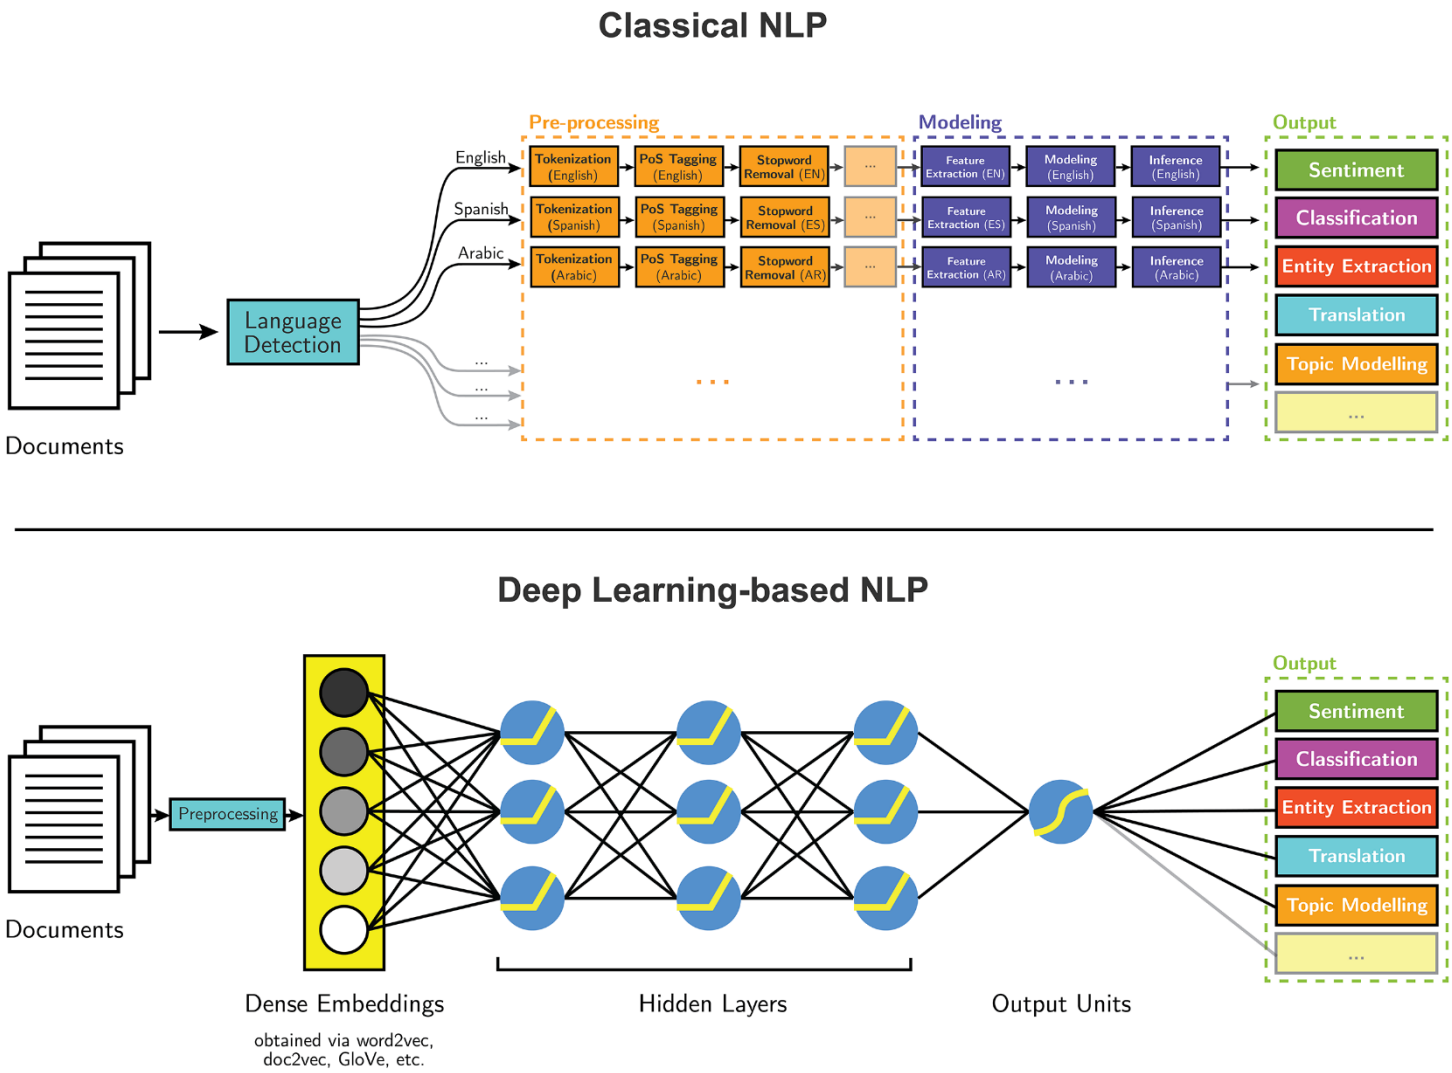

Source: Leveraging Deep Learning for Multilingual Sentiment Analysis - AYLIEN News API

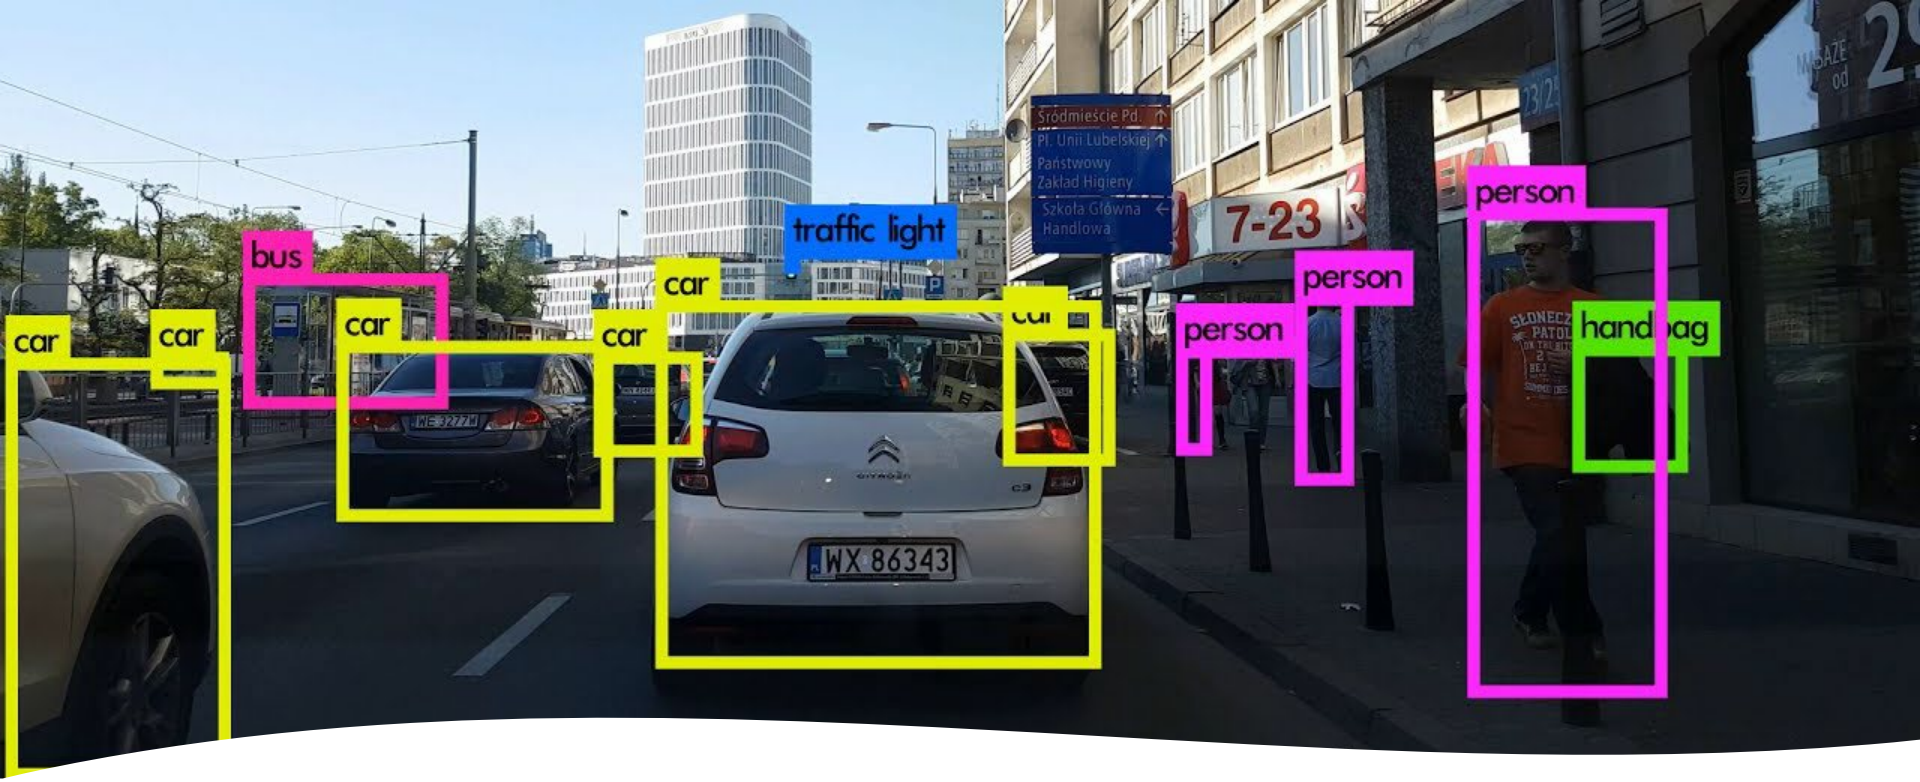

Computer Vision
Image Processing & Object Recognition

Uses Deep Convolutional Neural Networks

Robotics

- Hardware, sensors, control theory (feedback-based controllers)